[Alkemy]
Designed by humans. Powered by AI.

LUISS

AI Productivity analisys.

Beyond which AI usage threshold does rework destroy profit margin?

Presented by :

Group SEA:
· Simone Cavallaro (815441)
· Andrea Mammano (821311)
· Elisa Presciutti (812891)

Alkemy · Machine Learning Project
· LUISS A.Y. 2025/2026

[Alkemy]
Designed by humans. Powered by AI.

LUISS

Beyond which AI usage threshold does rework destroy profit margin?

AI makes tasks faster.

But faster does not always mean more profitable.

We asked: where does the gain go?

Is there a usage level beyond which **AI** starts **costing more than it saves**?

Is the damage done through **quality loss or** through **billable hours?**

Does the pricing structure determine who captures the **efficiency gain**?

[Alkemy]
Designed by humans. Powered by AI.

LUISS

AI improves margin: but the gains only consolidate above 50%

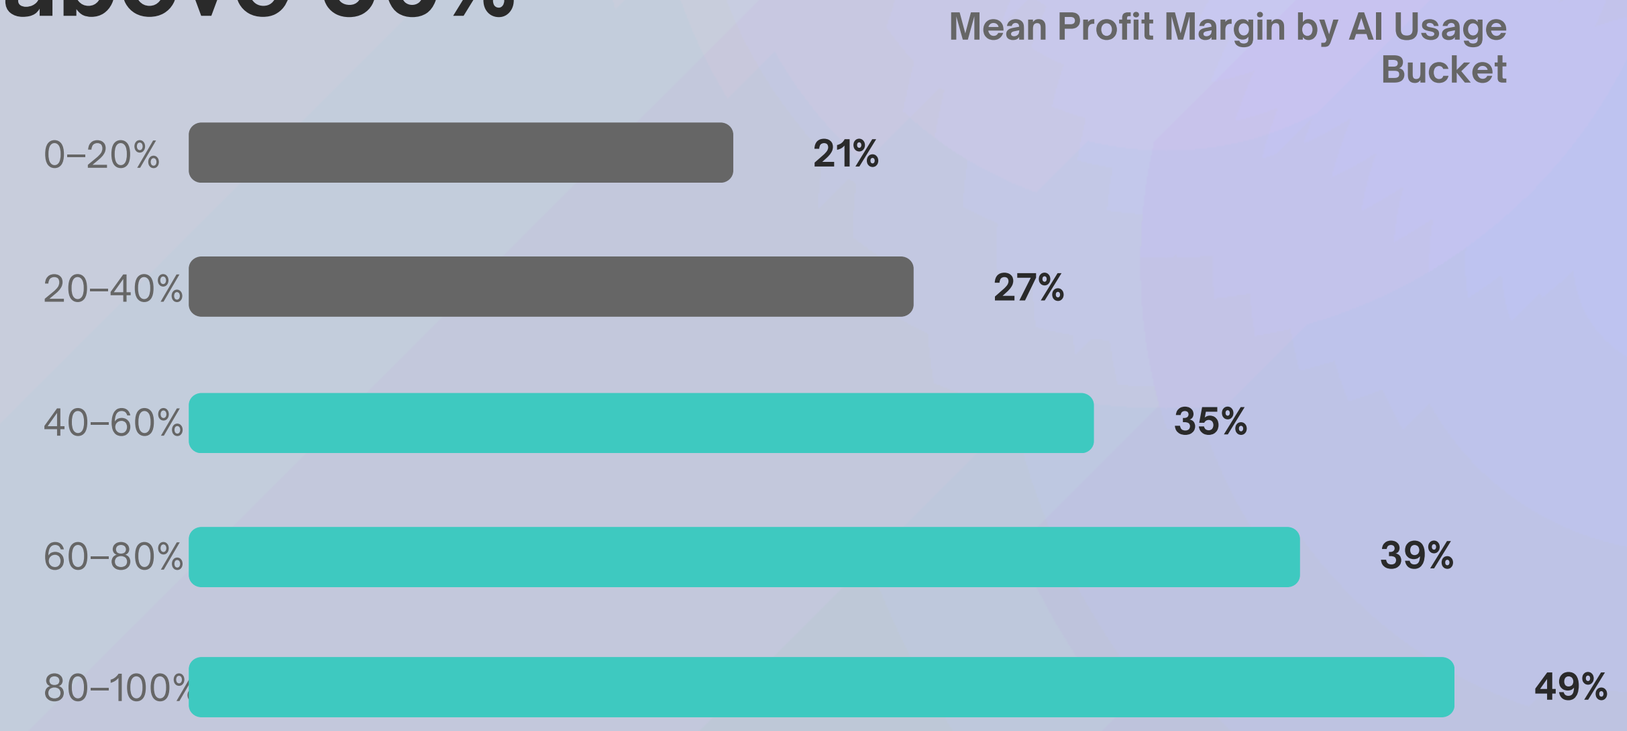

+28pp
margin gain from lowest to highest AI intensity bucket

−44%
SLA breach rate across the same range

31% → 19%
task loss rate falls as AI intensity rises

50–60%

Below this, rework absorbs the speed gain.

Above it, the efficiency dividend becomes real and measurable.

[Alkemy]
Designed by humans. Powered by AI.

LUISS

AI saves hours.
The contract decides who keeps them.

**Once hours_spent is controlled in OLS, all quality metrics become statistically insignificant ($p > 0.4$).
The economic mechanism is time, not quality.**

HOW AI ACTS ON PROFIT

The productivity gain is real.

On average, high-AI tasks save 2.0 hours of execution.
0.9 of those are absorbed by rework.
Net gain: 1.1 hours per task.

CONTRACT TYPE DECIDES OUTCOME

Same 1.1 hrs saved. Different P&L impact.

↑ Fixed / Value-based: hours saved = margin captured

↓ Hourly billing: hours saved = revenue lost

[Alkemy]
Designed by humans. Powered by AI.

LUISS

Under hourly contracts, AI makes Alkemy faster for the client. Not more profitable for itself.

Interaction term AI × hourly pricing: coef = −29.5 pp (p = 0.0005 ***) Margin under hourly peaks at 26%, then reverses

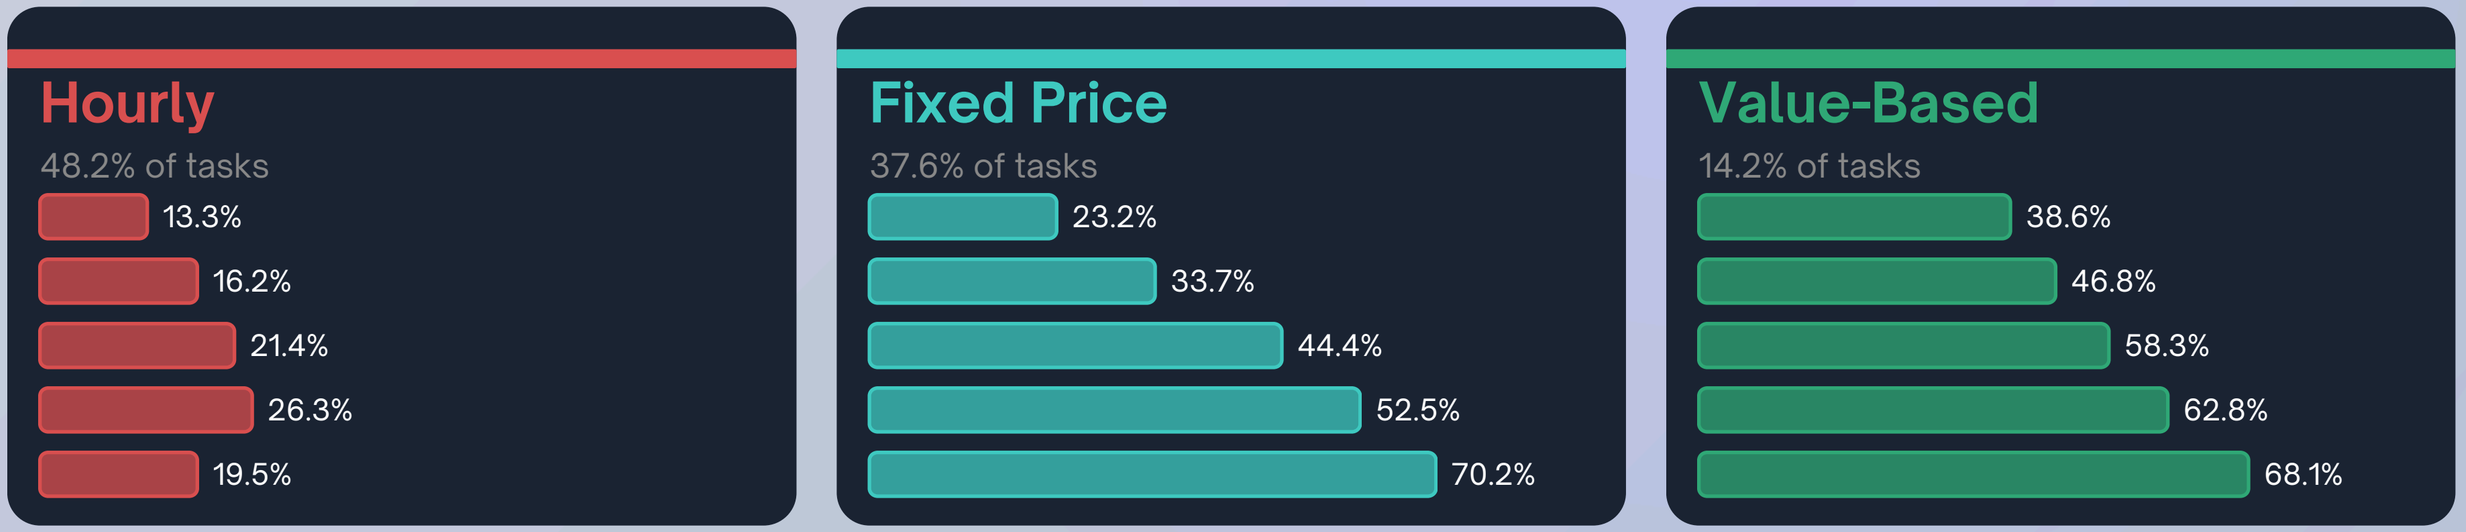

[Alkemy]
Designed by humans. Powered by AI.

LUISS

Task type determines how far you can push AI.

Task Type	Multiplier	Signal	Recommendation
Release	+4.8×		Push AI hard: biggest multiplier in the dataset
Ad	+2.1×		Strong response: cap at 80% to avoid rework spike
Dev	+2.0×		Consistent gains: works well at high intensity
Report	+1.6×		Good returns: push above 50%
Design	+1.5×		Moderate: quality stability at risk above 70%
Article	+0.8×		Cap at 40–60%: performance peaks in medium band

[Alkemy]
Designed by humans. Powered by AI.

LUISS

More rework, fewer losses: the paradox at the heart of AI adoption

THE PROFITABILITY PARADOX

Loss rate falls 31% → 19% while rework ratio rises to 32%. Speed creates an efficiency cushion that absorbs correction cost.

THE CRITICAL PIVOT: 50–60%

Between 50–60% AI usage the business crosses from high-loss to high-rework, low-risk. This is the operational break-even.

REWORK IS NOT FAILURE

The OLS quadratic term for rework is non-significant ($p = 0.605$). There is no statistical tipping point where rework destroys the business. Fast generation plus iteration is the profitable model.

Baron-Kenny mediation: 68% of AI's effect on profit margin flows through rework — not directly. Once rework is controlled, the direct AI coefficient becomes non-significant ($p = 0.645$). Fixing QA matters more than capping AI adoption.

[Alkemy]
Designed by humans. Powered by AI.

LUISS

THE THRESHOLD

Where the SHAP contribution of ai_usage_pct flips sign

30–60%

Below it, AI usage is on average a drag on predicted profit. Above it, it becomes a consistent positive contributor — but only if the pricing structure allows it.

BELOW 30%

−€52.5

average SHAP contribution per task in this zone

THE CROSSOVER

≈ €103

total swing in predicted profit attributable to AI alone

ABOVE 60%

+€51.1

average SHAP contribution per task in this zone

[Alkemy]
Designed by humans. Powered by AI.

LUISS

What actually drives profit and where AI sits in that picture (SHAP Analysis)

#1

PRICING MODEL — THE STRUCTURAL DRIVER

pricing_model_hourly ranks first in SHAP. Its impact on predicted profit is larger than ai_usage_pct across every segment of the data.

#2

BILLABLE RATIO — THE EFFICIENCY SIGNAL

How much of worked time becomes billed time is the second-strongest predictor. It captures whether the team runs lean or absorbs hidden slack.

#3

AI USAGE — 13TH OUT OF 57

ai_usage_pct has a real but indirect effect. It does not independently move profit — it amplifies whatever structure already exists around it.

1. pricing_model_hourly 15.3 2. billable_ratio 12.5 3. seniority_senior 9.6 4. hours_spent 7.8 · · · 13. ai_usage_pct 0.7 ← not a direct driver

[Alkemy]
Designed by humans. Powered by AI.

LUISS

MODEL VALIDATION

Three models, the same story

Model	Test R^2	MAE (pp)	CV R^2	Key insight
Linear Regression	0.205	32.1	0.083	Interpretable baseline
Lasso (LassoCV α = 0.57)	0.205	—	—	32 / 57 features carry signal
LR + Interaction (AI × hourly)	0.240	31.4	—	AI × hourly: p = 0.0005 ***
Random Forest (GridSearchCV)	**0.345**	27.6	0.184	Best predictive model

$\Delta R^2 = +0.14$

RF vs LR — strong non-linearity confirmed. Threshold effects are real, not artefacts.

59% of variance unexplained

Missing: client pricing power, AI tool quality, team capacity at the moment of assignment.

[Alkemy]
Designed by humans. Powered by AI.

LUISS

KEY INSIGHT

Three things that data taught us

01 AI is a multiplier, not a solution

AI usage does not independently move profit margin:

it amplifies whatever pricing structure already exists.

It changes between fixed and hourly contract.

The contract is the variable, not the AI.

02 The binary view was wrong

Tasks with any AI usage showed lower median margin than tasks with no AI.

This seems damaging: until you treat AI as intensity rather than a flag.

At 80–100% AI usage, median margin reaches 49%.

The binary comparison was obscuring the real signal entirely.

03 Rework rises with AI — and so does profitability

We expected rework to be the margin killer.

Instead, rework rate doubles as AI intensity rises (0.133 → 0.298) while margin rises from 21% to 49% over the same range.

The paradox forced us to rethink the mechanism: the intervention point is pricing, not quality control.

[Alkemy]
Designed by humans. Powered by AI.

LUISS

What we found, what we built, and what went wrong

Concrete Decision

The constraint is not the technology: it is the contract

AI can improve profitability, but only when adoption is deep enough and the commercial structure is aligned to capture the gains. Getting one without the other leaves value on the table. Nearly half the dataset runs on hourly contracts. That is the structural constraint. Reprice first, then scale.

Thing discovered thanks to AI

The rework paradox: a finding that would have been dismissed as noise

Combining visual and tabular analysis tools revealed a counterintuitive trend: higher AI intensity increases both rework costs and profit margins at the same time.
Rework is actually an indicator of task complexity, rather than just an isolated expense.
Without looking through these methods, this paradox would have likely been ignored instead of being understood.

Mistake made by AI

AI kept pushing us to cap the outliers. We did not. That was the right call.

Standard data cleanup methods would have removed valid extreme values by treating them as errors.
Accepting the AI's suggestion would have skewed the results toward the average, hiding important edge cases.
AI offered a theoretically solution but lacked business context. Catching the mistake required human judgment to understand what the outliers actually meant, rather than treating data preprocessing as a mindless checklist.

[Alkemy]
Designed by humans. Powered by AI.

LUISS

Three decisions that follow directly from the data

1
Primary lever

Reprice before scaling AI further

Increasing AI adoption under hourly contracts transfers the efficiency gain to the client, not the business. The pricing book needs to change first.

2
Immediate

Set a 30% floor, target 60%

Below 30%, AI usage is on average a net cost. The transition zone is 30–60%. Above 60%, profit and time savings consolidate. Treat these as operational thresholds, not guidelines.

3
Highest ROI

Rework reduction beats AI restriction

68% of AI's margin effect flows through rework. Improving QA and brief quality generates more margin uplift than capping AI usage rates.

[Alkemy]
Designed by humans. Powered by AI.

LUISS

BLIND SPOTS

What the data cannot tell us: yet.

The model is honest about its limits. These are the variables that would change the recommendations.

01

CLIENT RETENTION

We measure per-task profit. We do not know whether AI-heavy delivery affects renewal rates, NPS, or long-term contract value.

HIGHEST PRIORITY

02

TOOL QUALITY

ai_usage_pct tells us how much AI was used — not how well. Two tasks at 80% can have completely different quality profiles.

PROCUREMENT

03

REVISION CAUSALITY

Rework hours capture cost, not cause. We cannot distinguish AI-generated errors from scope changes or client preference shifts.

ROOT CAUSE

[Alkemy]
Designed by humans. Powered by AI.

LUISS

The AI is working. The pricing model isn't.

RESTRUCTURE THE CONTRACT BEFORE YOU SCALE THE TECHNOLOGY.

Group SEA — Simone Cavallaro · Andrea Mammano · Elisa Presciutti · LUISS 2025/2026